

Lejun Min

Researcher, Artist

- contact@lejun.site
- lejun.site
- LinkedIn
- GitHub

About me

Master's student at Stanford CCRMA specializing in music generation, audio signal processing, and human-computer interaction. Published at top venues including ICLR (Spotlight) and ISMIR.

Intermedia & sonic artist with various performances and installations. Passionate about bridging AI and creative expression through innovative music technologies.

Technical Skills

Research Areas

- Music Generation
- Music Information Retrieval
- Representation Learning
- Multimodal Learning
- Digital Signal Processing
- Human-Computer Interaction

Programming

- Python
- C++
- C
- Java
- Rust
- Verilog
- Lua
- Arduino
- WGSL
- Bash

Machine Learning

- PyTorch
- TensorFlow
- Accelerate
- Lightning


Audio & Graphics

- JUCE
- Chuck
- Reaper
- Pure Data
- Blender
- Adobe Premiere

Development

- Arch Linux
- Neovim
- Git
- Docker

Languages



Education

Stanford University, CCRMA

2024 – Present

M.A. in Music, Science, and Technology (GPA 4.0/4.0)

- Advisors: Prof. Marina Bosi, Prof. Takako Fujioka
- CCRMA Flagship Project Award, Chiang Chen Overseas Graduate Fellowship

Shanghai Jiao Tong University, Zhiyuan College

2019 – 2023

B.Eng. in Computer Science (Fellowship)

- Member of **ACM Honor Class** – elite CS program for top 5% students
- Top 1% Outstanding Bachelor's Thesis (41 out of 3873 graduates)

Experience

Smule Labs – Research Intern

Oct. 2025 – Dec. 2025, Mar. 2026 – May 2026

- Built a disentangled music semantic embedding by end-to-end audio captioning with self-supervised learning regularizations. Work under preparation for publication.
- Supervisor: Yongyi Zang

Sony CSL Paris – Research Intern

June 2025 – Sept. 2025

- Designed end-to-end mixing/mastering system using audio representation learning – pioneering study on automatic mixing with fully generative approach.
- Supervisor: Dr. Stefan Lattner

SoundPatrol & Stanford – Research Assistant

Feb. 2025 – May 2025

- Trained singer representation model for deepfake detection – reduced EER by 10%, achieved 96% top-1 accuracy on SingFake dataset.
- Advisor: Prof. John Thickstun, Prof. Walter De Brouwer

Music X Lab, MBZUAI – Research Assistant

June 2022 – Dec. 2022, Sept. 2023 – Feb. 2024

- Led development of hierarchical symbolic music generation (ICLR 2024) and polyphonic generation with diffusion models (ISMIR 2023).
- Advisor: Prof. Gus Xia

Publications

Leveraging Rotational M/S Coding and ML in Stereo Audio Coding

IWSSPA 2025

L. Min, S. Chen, M. Bosi

MuPT: A Generative Symbolic Music Pretrained Transformer

ICLR 2025

X. Qu, Y. Bai, ... L. Min, et al. [OpenReview] [Demo]

Whole-song Hierarchical Generation Using Cascaded Diffusion Models

ICLR 2024 (Spotlight)

Z. Wang, L. Min, G. Xia [OpenReview] [Demo]

Polyffusion: A Diffusion Model for Polyphonic Score Generation

ISMIR 2023

L. Min, J. Jiang, G. Xia, J. Zhao [Poster] [Demo]

Teaching

Perceptual Audio Coding, Teaching Assistant at Stanford

Winter 2026

Reinforcement Learning, Teaching Assistant at SJTU

Spring 2023

Design and Analysis of Algorithms, Teaching Assistant at SJTU

Spring 2022

Principle and Practice of Computer Algorithms, Teaching Assistant at SJTU

Summer 2021

Selected Projects

- Mx Compiler: Hand-made compiler surpassing -O1 optimization (Java)
- RISC-V CPU: 5-pipelined RISC-V32I with FPGA implementation (Verilog)
- Ray Tracer: Complete ray tracing engine in Rust
- A Chan Conversation: Ambisonics performance with Gametrak
- Umbrella: Intermedia piece for Ambisonics audio & video exploring self and fear
- Sonic Skateboard: Arduino + Chuck musical instrument